




SLO goodput (requests/s) - higher is better



Output tokens/s - higher is better



p99 first-token wait (s) - lower is better



GPU kernel-busy time (%) - diagnostic

